



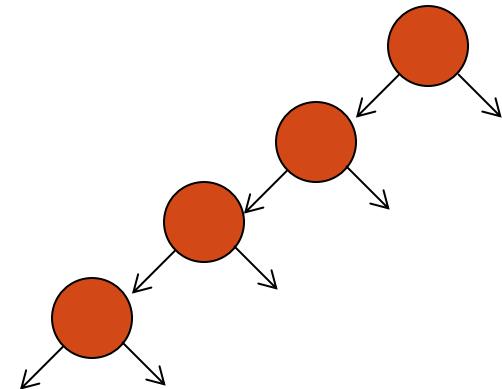
PRIORITY QUEUES

PRIORITIZATION PROBLEMS

- Printer Jobs : Print jobs print in a queue, what if we want the faculty jobs to print before the student jobs, irrespective of which order they are sent to the printer.
- ER Scheduling : Treating a gunshot patient before a guy with a sore neck irrespective of the time of their arrival.
- *Why can't we solve these problems efficiently with the data structures we have (list, sorted list, queue, stack etc.)?*

PROBLEMS WITH EXISTING DATA STRUCTURES

- *list* : store jobs in a list; remove min/max by searching ($O(N)$)
 - problem: expensive to search
- *sorted list* : store in sorted list; binary search it in $O(\log N)$ time
 - problem: expensive to add/remove ($O(N)$)
- *binary search tree* : store in BST, go right for max in $O(\log N)$
 - problem: tree becomes unbalanced



PRIORITY QUEUE

- **Priority Queue:** a collection of ordered elements that provides fast access to the minimum (or maximum) element
 - usually implemented using a tree structure called a *heap*
- Each element of the queue has some priority, and all elements are processed based on their priorities
- An element with the high priority is processed before element with low priority.
- FIFO rule applies if two elements have same priority
- Implement it as a sorted linked list, sorted on priority
- Element with highest priority is always inserted first

PRIORITY QUEUE ADT

A priority queue P supports the following methods:

- size(): Return the number of elements in P
- isEmpty(): Test whether P is empty
- insertItem(k,e): Insert a new element e with key k into P
- minElement(): Return (but don't remove) an element of P with smallest key; an error occurs if P is empty.
- minKey(): Return the smallest key in P; an error occurs if P is empty
- removeMin(): Remove from P and return an element with the smallest key; an error conditions occurs if P is empty.

LIST BASED PRIORITY QUEUE

- Unsorted list implementation

- Store the items of the priority queue in a list-based sequence, in arbitrary order



- Performance:

- `insertItem` takes $O(1)$ time since we can insert the item at the beginning or end of the sequence
- `removeMin`, `minKey` and `minElement` take $O(n)$ time since we have to traverse the entire sequence to find the smallest key

- Sorted list implementation

- Implementation

- Store the items of the priority queue in a sequence, sorted by key



- Performance:

- `insertItem` takes $O(n)$ time since we have to find the place where to insert the item
- `removeMin`, `minKey` and `minElement` take $O(1)$ time since the smallest key is at the beginning of the sequence

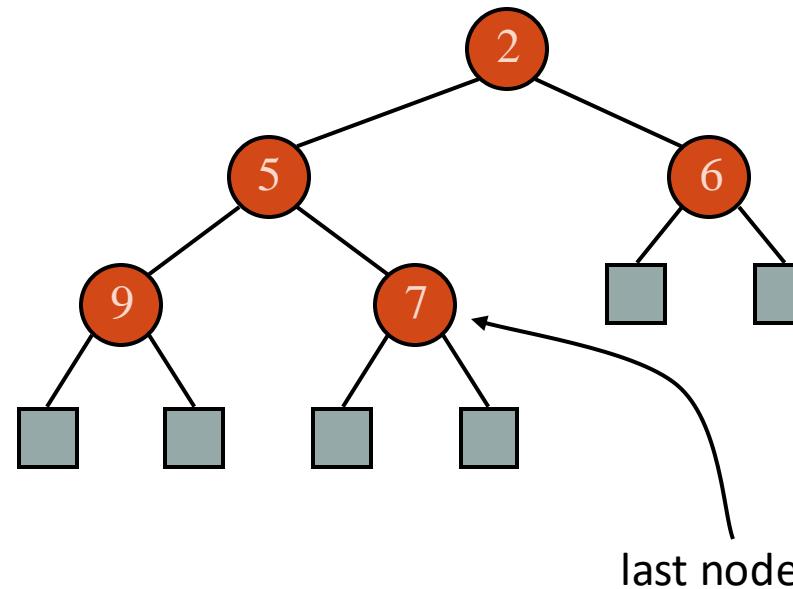
7

IMPLEMENTATION OF A PRIORITY QUEUE

HEAPS

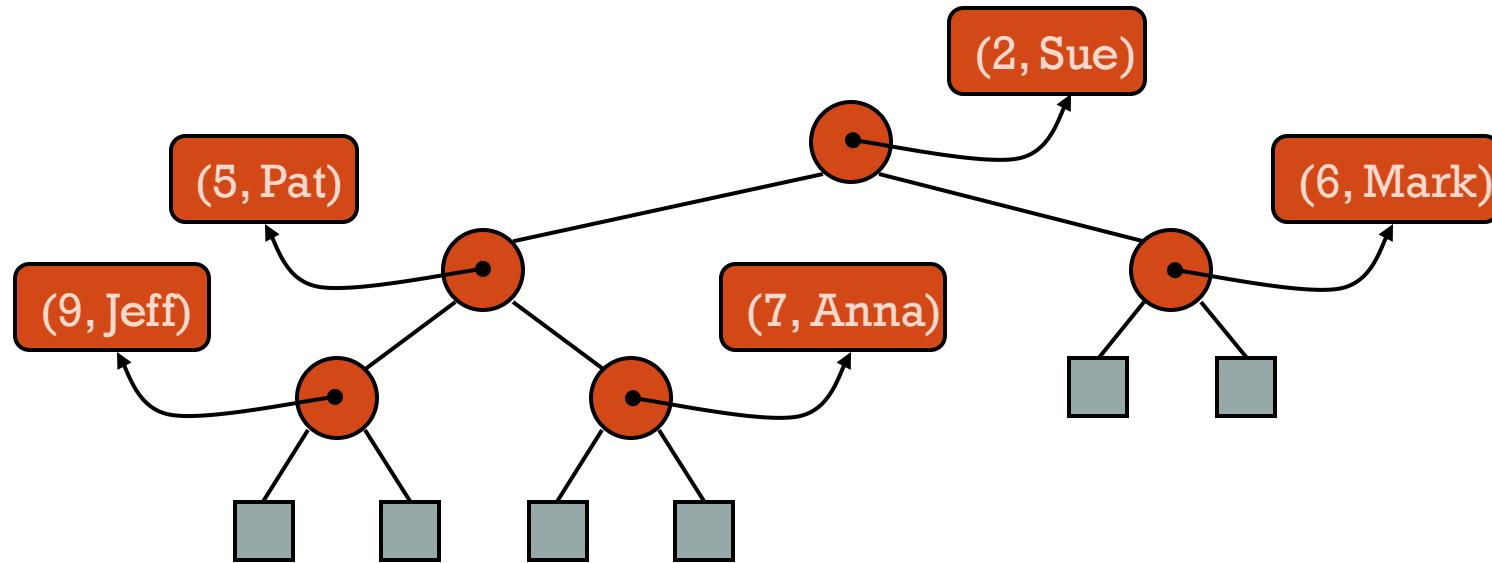
- A heap is a binary tree storing keys at its internal nodes and satisfying the following properties:
 - **Heap-Order:** for every internal node v other than the root, $\text{key}(v) \geq \text{key}(\text{parent}(v))$
 - **Complete Binary Tree:** let h be the height of the heap
 - for $i = 0, \dots, h - 1$, there are 2^i nodes of depth i
 - at depth $h - 1$, the internal nodes are to the left of the leaf nodes

- The last node of a heap is the rightmost internal node of depth $h - 1$



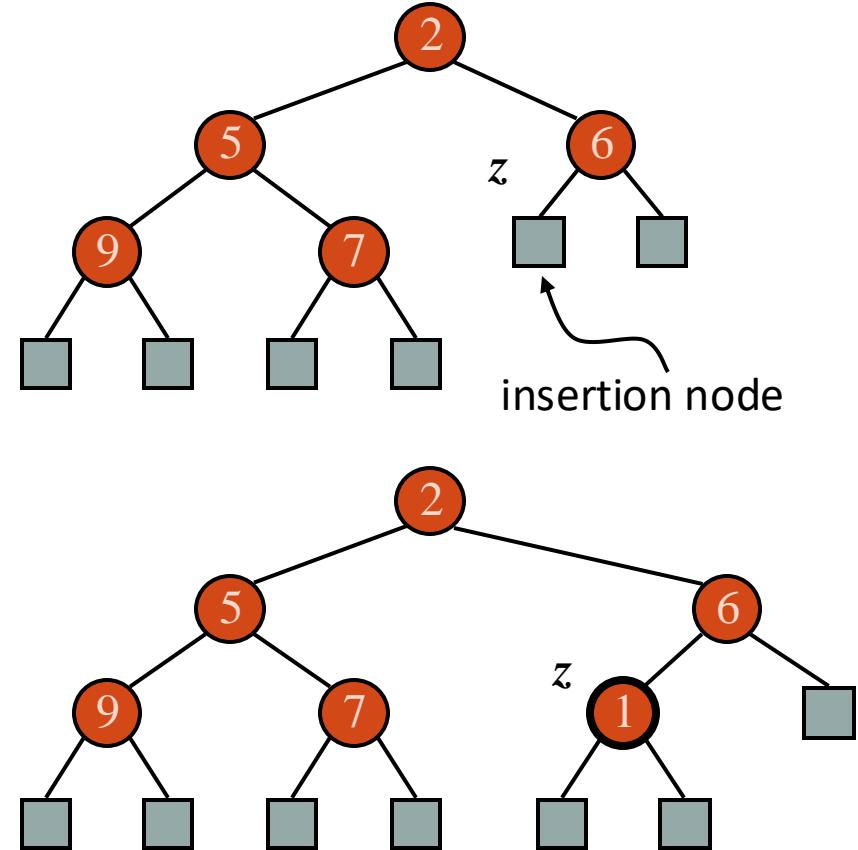
HEAPS AND PRIORITY QUEUES

- We can use a heap to implement a priority queue
- We store a (key, element) item at each internal node
- We keep track of the position of the last node
- For simplicity, we show only the keys in the pictures



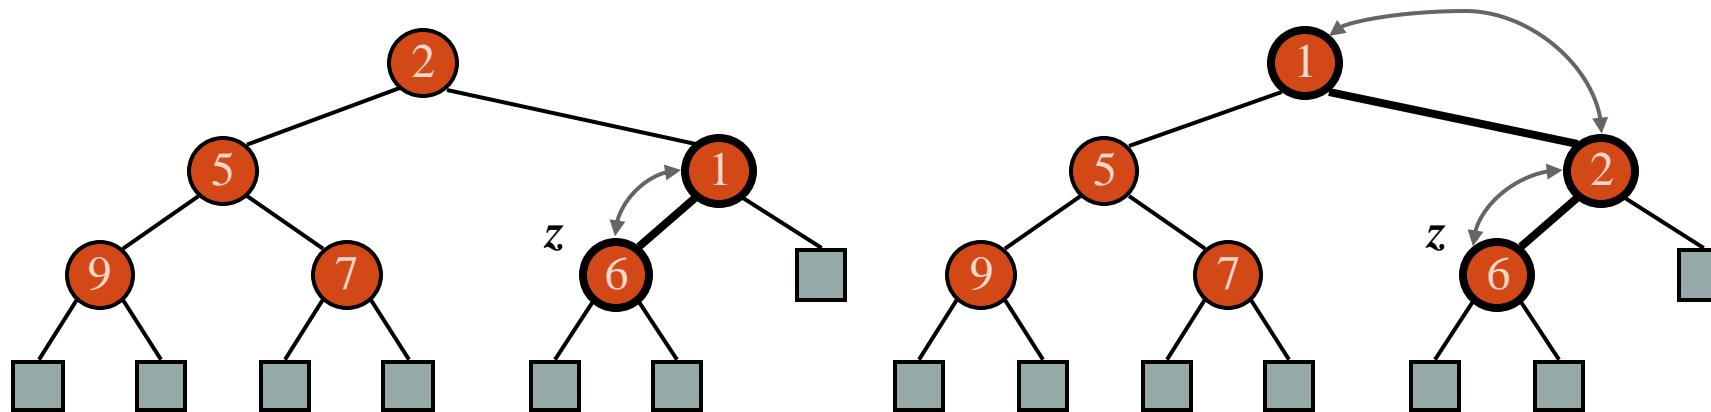
INSERTION INTO A HEAP

- Method `insertItem` of the priority queue ADT corresponds to the insertion of a key k to the heap
- The insertion algorithm consists of three steps
 - Find the insertion node z (the new last node)
 - Store k at z and expand z into an internal node
 - Restore the heap-order property (discussed next)



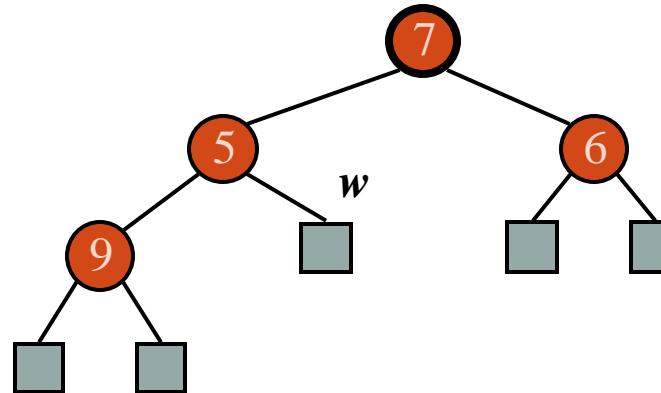
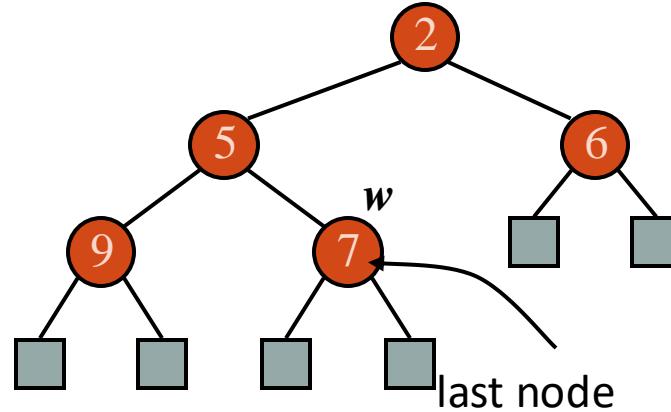
UPHEAP

- After the insertion of a new key k , the heap-order property may be violated
- Algorithm upheap restores the heap-order property by swapping k along an upward path from the insertion node
- Upheap terminates when the key k reaches the root or a node whose parent has a key smaller than or equal to k
- Since a heap has height $O(\log n)$, upheap runs in $O(\log n)$ time



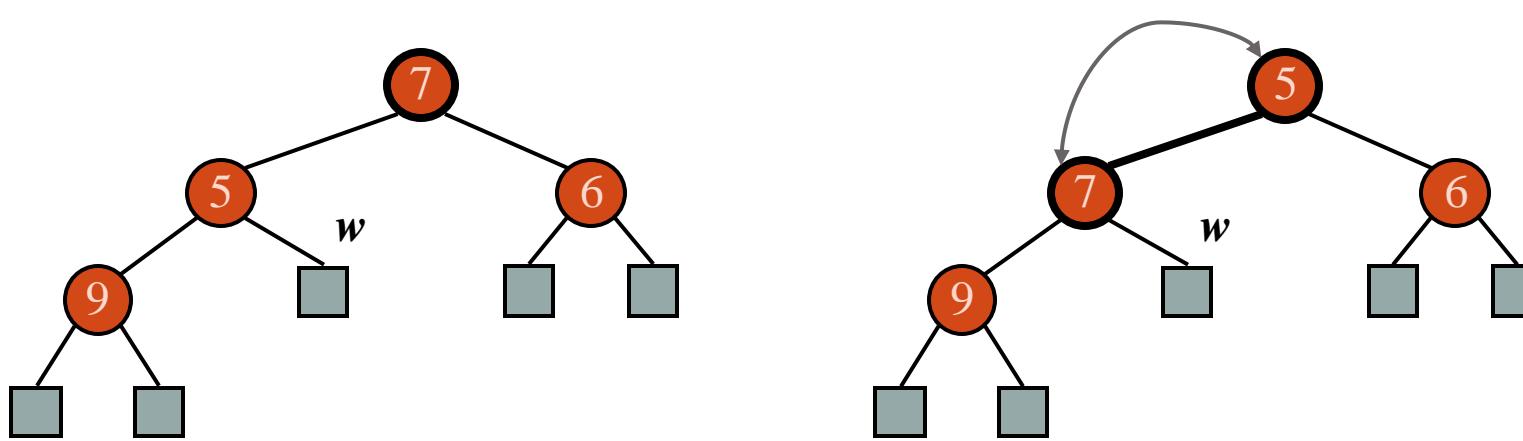
REMOVAL FROM A HEAP

- Method `removeMin` of the priority queue ADT corresponds to the removal of the root key from the heap
- The removal algorithm consists of three steps
 - Replace the root key with the key of the last node w
 - Compress w and its children into a leaf
 - Restore the heap-order property (discussed next)



DOWNHEAP

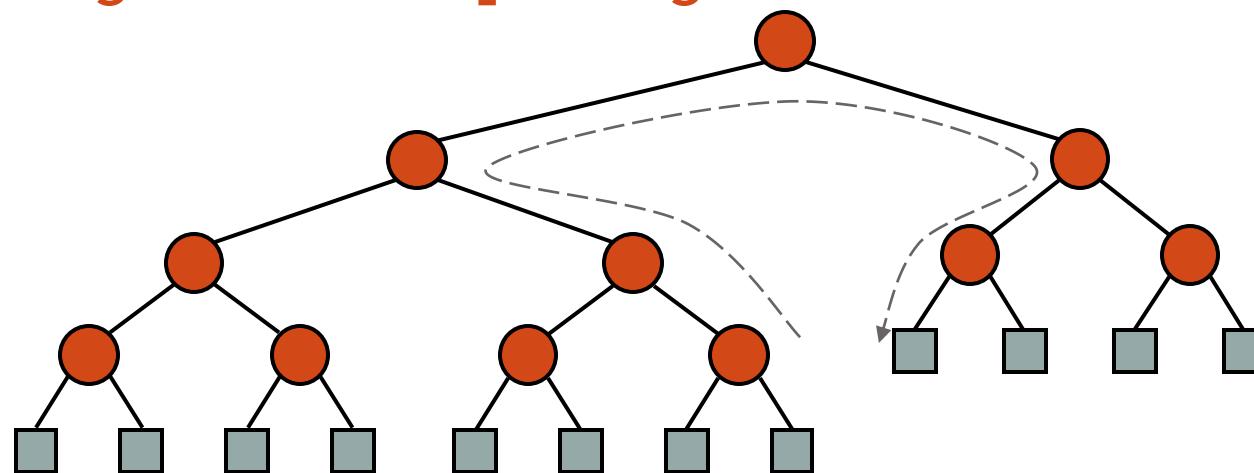
- After replacing the root key with the key k of the last node, the heap-order property may be violated
- Algorithm downheap restores the heap property by swapping key k with the child with the smallest key along a downward path from the root
- Downheap terminates when key k reaches a leaf or a node whose children have keys greater than or equal to k
- Since a heap has height $O(\log n)$, downheap runs in $O(\log n)$ time



- Start with the parent
- Compare it with both its left child and right child
 - If both children are greater than k – *Nothing to be done*
 - If one child is smaller than k – *the Smaller child value is moved up*
 - If both children are smaller than k – *the Smaller of the two children is moved up.*
- Stop when both children are greater than k or we reach a leaf node.

UPDATING THE LAST NODE

- The insertion node can be found by traversing a path of $O(\log n)$ nodes
 - Go up until a left child or the root is reached
 - If a left child is reached, go to the right child
 - Go down left until a leaf is reached
- Similar algorithm for updating the last node after a removal



HEAP SORT – A PEEK

- Consider a priority queue with n items implemented by means of a heap
 - the space used is $O(n)$
 - methods `insertItem` and `removeMin` take $O(\log n)$ time
 - methods `size`, `isEmpty`, `minKey`, and `minElement` take time $O(1)$ time
- Using a heap-based priority queue, we can sort a sequence of n elements in $O(n \log n)$ time
- The resulting algorithm is called heap-sort
- Heap-sort is much faster than quadratic sorting algorithms, such as insertion-sort and selection-sort

REPRESENTING HEAPS - SEQUENTIALLY

- We can represent a heap with n keys by means of an array of length n
- For the node at rank i
 - the left child is at rank $2i$ (*if $2i > n$ no left child*)
 - the right child is at rank $2i + 1$ (*if $2i+1 > n$ no right child*)
 - parent is $i/2$
- Links between nodes are not explicitly stored
- The leaves are not represented

